

Multi-Class Text Classification: A Comparison of Word Representations and ML/NN Models

Ishraq Kamal Adib

Dept. of

Computer Science and Engineering
BRAC University
Dhaka, Bangladesh
ishraq.kamal.adib@g.bracu.ac.bd

Mostakim Morshed

Dept. of

Computer Science and Engineering
BRAC University
Dhaka, Bangladesh
mostakim.morshed@g.bracu.ac.bd

Stanley Matthew Das

Dept. of

Computer Science and Engineering
BRAC University
Dhaka, Bangladesh
stanley.matthew.das@g.bracu.ac.bd

Abstract—This project applies a comprehensive experimental pipeline to a multi-class question-answer topic classification dataset. We performed exploratory data analysis (EDA), applied standard text preprocessing, implemented four word representations (Bag-of-Words, TF-IDF, GloVe, Skip-gram), and trained a suite of models: three classic machine learning classifiers (Logistic Regression, Naive Bayes, Random Forest) and seven neural network architectures (DNN, SimpleRNN, GRU, LSTM, and their bidirectional variants). We trained the ML models on BoW, TF-IDF; all NN architectures on GloVe, Skip-gram with and DNN on all four. The best ML model found was Logistic Regression with TF-IDF and the best NN model was a Bidirectional GRU with Skip-gram embeddings.

Index Terms—text classification, TF-IDF, BoW, GloVe, Skip-gram, RNN, LSTM, GRU, word embeddings

I. INTRODUCTION

This project addresses multi-class classification of short question titles into topical labels (e.g. Science & Mathematics, Education & Reference, Politics & Government, Entertainment & Music, Sports, etc.). The goal is to compare classical feature-based classifiers with neural sequence models across different word representations and to document reproducible experimental choices and outcomes. The provided dataset was already split into an 80% training set and 20% test set; 10% of the training set was further used for validation of the training data during hyperparameter tuning and early stopping.

II. METHODOLOGY

A. Dataset & Exploratory Data Analysis

We began by examining the dataset to understand its structure and distribution before designing preprocessing and modeling strategies. The dataset contains short question titles annotated with multiple topical categories. To characterize the data, we performed exploratory data analysis (EDA) and observed the following:

- **Word Cloud:** Frequently occurring words include “question,” “answer,” “content,” and “title,” along with high-frequency verbs such as “say,” “know,” and “want.” This suggests the dataset contains both domain-specific and conversational vocabulary.
- **Class Distribution:** Classes are relatively balanced, with no extreme skew, although certain categories such as

Education & Reference and Science & Mathematics are slightly more dominant.

- **Text Lengths:** Most samples are very short, typically below 20 tokens, reinforcing the short-text nature of the task.
- **Missing Values:** The dataset has no missing values across relevant fields (QA text, class labels, length/count features).




Fig. 1. EDA visualizations: (a) class distribution, (b) text length distribution, (c) word count distribution, (d) text length by class, (e) word count by class, and (f) missing values heatmap.

These EDA results provided the basis for preprocessing decisions. For example, we retained topical nouns and key tokens while removing URLs and user handles, since they were identified as noise. Furthermore, the short length of most titles suggested that recurrent architectures would not require deep stacking to capture dependencies effectively.

B. Preprocessing

The preprocessing pipeline implemented in the notebook includes:

- Lowercasing, trimming whitespace.
- punctuation stripping except where meaningful.
- Tokenization, stopword removal, and lemmatization



Fig. 2. Word cloud of the training dataset showing most frequent tokens. Common topical words and conversational markers dominate

- For embedding experiments, an explicit OOV embedding initialization was used for words.
 - After visualizing wordcloud of the above figure we preprocessed some unnecessary words which were very frequent

C. Word Representations

We used four representations required by the project:

- Bag-of-Words (BoW): scikit-learn CountVectorizer (unigram; optional n-grams in experiments), tuned `min_df` and `max_features`.
 - TF-IDF: scikit-learn TfidfVectorizer with `sublinear_tf`, tuned `ngram_range` and `max_df/min_df`.
 - GloVe: pretrained GloVe embeddings loaded and mapped to the dataset vocabulary.
 - Skip-gram: skip-gram embeddings trained on the training corpus (gensim/Keras implementation) and used to initialize the embedding layer for sequence models.

D. Model Architectures

Machine Learning Models:

- **Logistic Regression:** A linear model used with both Bag-of-Words (BoW) and TF-IDF representations to classify text by finding a linear decision boundary.

- **Hyperparameters:** max_iter=5000

- **Naive Bayes (MultinomialNB):** A probabilistic model suitable for text classification with discrete features (like word counts in BoW or TF-IDF). It's based on applying Bayes' theorem with the “naive” assumption of conditional independence between features.

- **Hyperparameters:** Default parameters were used.

- **Random Forest:** An ensemble learning method that builds multiple decision trees and merges their predictions to improve accuracy and control overfitting. It was applied with both BoW and TF-IDF features.

- **Hyperparameters:** n_estimators=100,
random_state=42

Deep Learning Models:

- **Deep Neural Network (DNN):** A feedforward neural network with multiple layers using learning rate of 0.001. Used with BoW, TF-IDF, and averaged word embeddings (Skip-gram and GloVe) to learn non-linear relationships in the data.

- Hyperparameters:

- * Hidden Layers: 2
 - * Neurons per layer: 256 (first), 128 (second)
 - * Activation function: ReLU
 - * Dropout: 0.5 (after first), 0.3 (after second)
 - * Output layer activation: Softmax
 - * Optimizer: Adam
 - * Loss function: Sparse Categorical Crossentropy
 - * Epochs: 10
 - * Batch Size: 32

- **Recurrent Neural Networks (RNNs):** Used with pre-trained GloVe and Skip-gram embeddings (embedding dimension 100), with embedding layer trainable=False.

- Hyperparameters (Common to SimpleRNN, GRU, and LSTM):

- * RNN layer units: 128
 - * Dropout: 0.2 (within RNN layer), 0.3 (after RNN layer), 0.2 (after first Dense)
 - * Recurrent Dropout: 0.2
 - * Dense layer neurons: 128
 - * Activation: ReLU (Dense layers)
 - * Output activation: Softmax
 - * Optimizer: Adam
 - * Loss: Sparse Categorical Crossentropy
 - * Epochs: 5
 - * Batch Size: 256

- **Bidirectional RNNs (BiRNNs):** SimpleRNN, GRU, or LSTM layer wrapped in a Bidirectional layer.

4. Hyperparameter Tuning and Validation Protocol

Manual tuning guided by validation (held-out from training) prioritized weighted F1 as the model selection metric. We changed number of epochs to tune.

III. RESULTS

This section reports the final test set metrics extracted from the notebook outputs and provides analysis and comparison.

A. Quantitative Results

Best models:

- a) Best ML (vector models): Logistic Regression (TF-IDF) — Test Acc 0.6290, F1 0.6248.:
 - b) Best NN (sequence models): Bidirectional GRU (Skip-gram) — Test Acc 0.7111, F1 0.7047.:. Worst models:
 - c) Worst ML (vector models): Random Forest (BoW) — Test Acc 0.5956, F1 0.5910.:

TABLE I
MODEL PERFORMANCE: TEST ACCURACY




Fig. 3. Accuracy Comparison of All 22 Combinations




Fig. 4. F1 Score Comparison of All 22 Combinations




Fig. 5. Example of a figure caption.

d) Worst NN (sequence models): Simple RNN for both combination — Test Acc 0.1244, F1 0.703.:

B. Comparative Analysis

a) Representation impact: BoW/TF-IDF + linear models provide robust baselines ($\approx 0.62\text{--}0.65$). Sequence models with GloVe or Skip-gram embeddings outperform these baselines when architectures are properly configured — particularly BiGRU/BiLSTM.:

b) Architecture impact: GRU and LSTM (especially bidirectional) consistently beat SimpleRNN, indicating vanishing




Fig. 6. Example of a figure caption.

gradient mitigation and better memory of longer context help, even for short titles.:

c) Stability issues: Some pairings (e.g., DNN with Skip-gram, SimpleRNN with embeddings) produced very low accuracies. Likely causes include poor embedding initialization, unsuitable learning rates, or architecture-representation mismatch.:

C. Confusion Matrices & Per-class Performance

Full confusion matrices and classification reports for top models are included in the notebook outputs. They show that most confusion occurs between topically adjacent classes (e.g., politics vs society). If desired, we can embed the notebook's confusion matrix images into the final PDF.

IV. CONCLUSION

We executed the full required experiment matrix and identified clear trends: for this short-text topic classification task, bidirectional recurrent models with Skip-gram embeddings yielded the best performance (BiGRU: Acc = 0.7111). TF-IDF with Logistic Regression give strong, fast baselines (Acc ≈ 0.6474). Compared to other NN models, DNN model shows disappointing results.

Limitations: Training NN models on such large datasets are computationally expensive on small hardware.

Future work: Future works can evaluate transformer-based encoders (BERT) and compare compute/accuracy tradeoffs.

REFERENCES

- [1] Seaborn Seaborn: statistical data visualization. Version 0.13.2. Online documentation. Available: [seaborn.pydata.org](https://seaborn.pydata.org/seaborn.pydata.org)
- [2] Pandas pandas documentation. Version 2.3.1, July 07, 2025. Online documentation. Available: [pandas.pydata.org](https://pandas.pydata.org/Pandas)
- [3] NLTK (Natural Language Toolkit) Natural Language Toolkit. Online platform. Available: nltk.org/nltk.org
- [4] wordcloud wordcloud · PyPI (Python Package Index). Latest version documentation. Available: pypi.org
- [5] M. Young, The Technical Writer's Handbook. Mill Valley, CA: University Science, 1989.